

You can now deploy any containerized application to your cluster. To keep things familiar, let's deploy Nginx using *Deployments* and *Services* to see how this application can be deployed to the cluster. You can use the commands below for other containerized applications as well, provided you change the Docker image name and any relevant flags (such as ports and volumes).

Still within the master node, execute the following command to create a deployment named nginx:

```
kubernetes-master:~$kubectl create deployment nginx --image=nginx
```

```
ubuntu@master-node: ~  
ubuntu@master-node:~$ kubectl create deployment nginx --image=nginx  
deployment.apps/nginx created  
ubuntu@master-node:~$ kubectl expose deploy nginx --port 80 --target-port 80 --type NodePort  
service/nginx exposed  
ubuntu@master-node:~$ kubectl get services  
NAME          TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE  
httpenv       ClusterIP     10.100.5.228   <none>         8888/TCP         3m8s  
kubernetes    ClusterIP     10.96.0.1     <none>         443/TCP          12m  
nginx         NodePort      10.110.112.52 <none>         80:32299/TCP     8s  
ubuntu@master-node:~$
```

A deployment is a type of Kubernetes object that ensures there's always a specified number of pods running based on a defined template, even if the pod crashes during the cluster's lifetime. The above deployment will create a pod with one container from the Docker registry's [Nginx Docker Image](#).

Next, run the following command to create a service named nginx that will expose the app publicly. It will do so through a *NodePort*, a scheme that will make the pod accessible through an arbitrary port opened on each node of the cluster:

```
kubernetes-master:~$kubectl expose deploy nginx --port 80 --target-port 80 --type NodePort
```

```
ubuntu@master-node:~$ kubectl expose deploy nginx --port 80 --target-port 80 --type NodePort  
service/nginx exposed
```

Services are another type of Kubernetes object that expose cluster internal services to clients, both internal and external. They are also capable of load balancing requests to multiple pods, and are an integral component in Kubernetes, frequently interacting with other components.

Run the following command:

```
kubernetes-master:~$kubectl get services
```

```
ubuntu@master-node:~$ kubectl get services
NAME         TYPE        CLUSTER-IP    EXTERNAL-IP  PORT(S)    AGE
httpenv      ClusterIP   10.100.5.228  <none>       8888/TCP   3m8s
kubernetes   ClusterIP   10.96.0.1     <none>       443/TCP    12m
nginx        NodePort    10.110.112.52 <none>       80:32299/TCP 8s
```

This will output text similar to the following:

Output

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
------	------	------------	-------------	---------	-----

kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	1d
------------	-----------	-----------	--------	---------	----

nginx	NodePort	10.109.228.209	<none>	80:nginx_port/TCP	40m
-------	----------	----------------	--------	-------------------	-----

From the third line of the above output, you can retrieve the port that Nginx is running on. Kubernetes will assign a random port that is greater than 30000 automatically, while ensuring that the port is not already bound by another service.

Note: if you're running your setup on ec2 ensure the nginx_port is open under the inbound rules in the security groups.

To test that everything is working, visit

http://worker_1_ip:nginx_port or

http://worker_2_ip:nginx_port

through a browser on your local machine. You will see Nginx's familiar welcome page. To see the deployed container on the worker node switch to worker01

```
on-slave#docker ps
```

Output: you will see the container for nginx image running.

If you want to scale up the replicas for a deployment (nginx in our case) the use the following command:

```
kubernetes-master:~$kubectl scale --current-replicas=1 --replicas=2  
deployment/nginx
```

```
kubernetes-master:~$kubectl get pods
```

Output: you will see 2/2 as output in nginx deployment.

```
ubuntu@master-node:~$ kubectl scale --current-replicas=1 --replicas=2 deployment/nginx  
deployment.apps/nginx scaled  
ubuntu@master-node:~$ kubectl get pods  
NAME                                READY   STATUS              RESTARTS   AGE  
httpenv-7b94c48ff5-djc9x           0/1     ImagePullBackOff    0           10m  
httpenv-7b94c48ff5-qkw6g           0/1     ImagePullBackOff    0           10m  
httpenv-7b94c48ff5-vvbbg           0/1     ImagePullBackOff    0           10m  
nginx-76d6c9b8c-jxlnx              1/1     Running             0           7s  
nginx-76d6c9b8c-pxdq9              1/1     Running             0           7m22s
```

```
kubernetes-master:~$kubectl describe deployment/nginx
```

Output: give details about the service deployed

```

ubuntu@master-node: ~
ubuntu@master-node:~$ kubectl describe deployment/nginx
Name: nginx
Namespace: default
CreationTimestamp: Tue, 20 Sep 2022 10:37:35 +0000
Labels: app=nginx
Annotations: deployment.kubernetes.io/revision: 1
Selector: app=nginx
Replicas: 2 desired | 2 updated | 2 total | 2 available | 0 unavailable
StrategyType: RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels: app=nginx
  Containers:
    nginx:
      Image: nginx
      Port: <none>
      Host Port: <none>
      Environment: <none>
      Mounts: <none>
      Volumes: <none>
Conditions:
  Type           Status  Reason
  ----           -
  Progressing    True    NewReplicaSetAvailable
  Available      True    MinimumReplicasAvailable
OldReplicaSets: <none>
NewReplicaSet: nginx-76d6c9b8c (2/2 replicas created)
Events:
  Type      Reason      Age   From      Message
  ----      -

```

If you would like to remove the Nginx application, first delete the nginx service from the master node:

```
kubernetes-master:~$ kubectl delete service nginx
```

```

ubuntu@master-node:~$ kubectl delete service nginx
service "nginx" deleted

```

Run the following to ensure that the service has been deleted:

```
kubernetes-master:~$ kubectl get services
```

You will see the following output:

```

ubuntu@master-node:~$ kubectl get services
NAME          TYPE          CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
httpenv       ClusterIP     10.100.5.228    <none>           8888/TCP         10m
kubernetes    ClusterIP     10.96.0.1       <none>           443/TCP          20m

```

Output

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
------	------	------------	-------------	---------	-----

kubernetes ClusterIP 10.96.0.1 <none> 443/TCP 1d

Then delete the deployment:

```
kubernetes-master:~$kubectl delete deployment nginx
```

```
ubuntu@master-node:~$ kubectl delete deployment nginx
deployment.apps "nginx" deleted
```

Run the following to confirm that this worked:

```
kubernetes-master:~$kubectl get deployments
```

```
deployment.apps "nginx" deleted
ubuntu@master-node:~$ kubectl get deployments
NAME          READY    UP-TO-DATE    AVAILABLE    AGE
httpenv       0/3      3             0            11m
ubuntu@master-node:~$
```

Output

No resources found.